

# LAB 10

Try CNN on "Fruit" dataset. Also modify number of layers and observe the performance difference:

<https://www.kaggle.com/moltean/fruits> (<https://www.kaggle.com/moltean/fruits>)

Or (In a case if you don't have that much dataPack available, download 20 images of apple and 20 images of orange from the internet and work on it with RANDOM state=Rollnumber stratergy, 80-20% training-testing division)

```
In [2]: # Import libraries
import numpy as np
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
from torch.autograd import Variable
from torch.utils.data import DataLoader
from sklearn.datasets import load_files
from sklearn.model_selection import train_test_split
```

```
In [3]: # Load Data Directory
data_dir = '/home/nihar/Desktop/SEM 7/ML/Lab/Lab9/sample-fruits-360'
```

```
In [4]: # Function for load images
def load_dataset(path):
    data = load_files(path)
    files = np.array(data['filenames'])
    targets = np.array(data['target'])
    target_labels = np.array(data['target_names'])
    return files, targets, target_labels
```

```
In [5]: # Load Dataset
x, y, target_labels = load_dataset(data_dir)
print("Dataset Loaded !")

# Get Training size and Test size
print('Total set size : ',x.shape)
print('Total targets : ',len(target_labels) )
```


```
Dataset Loaded !
Total set size : (975,)
Total targets : 65
```

```
In [6]: # Function for convert image to array
def convert_image_to_array(files):
    images_as_array=[]
    for file in files:
        images_as_array.append(keras.preprocessing.image.img_to_array(keras.
preprocessing.image.load_img(file)))
    return images_as_array

# Convert images to numpy array using keras.preprocessing library
x = np.array(convert_image_to_array(x),np.float32)
print(x.shape)

(975, 100, 100, 3)
```

```
In [7]: # Plot image on random data
plt.imshow(x[1]/255)
plt.title(target_labels[y[1]])
plt.show()
```



```
In [8]: # Flatten the features of image
x = x.reshape([-1,100*100*3])
x = x/255
print("final shape : " , x.shape)

final shape : (975, 30000)
```

```
In [9]: # Train and Test split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_st
ate=129)

# Get size of all set
print("X Train size : ", X_train.shape)
print("X Test size : ", X_test.shape)
print("Y Train size : ", y_train.shape)
print("Y Test size : ", y_test.shape)

X Train size : (780, 30000)
X Test size : (195, 30000)
Y Train size : (780,)
Y Test size : (195,)
```

```
In [10]: # Convert numpy array to torch
X_train = torch.from_numpy(X_train)
y_train = torch.from_numpy(y_train).type(torch.LongTensor)
X_test = torch.from_numpy(X_test)
y_test = torch.from_numpy(y_test).type(torch.LongTensor)
```

```
In [11]: # Define no of iteration, batch size, num_epochs
batch_size=100
n_iters = 1000
num_epochs = n_iters / (len(X_train) / batch_size)
num_epochs = int(num_epochs)
```

```
In [12]: # Set train and test
train = torch.utils.data.TensorDataset(X_train,y_train)
test = torch.utils.data.TensorDataset(X_test,y_test)
train_loader = DataLoader(train, batch_size = batch_size, shuffle = False)
test_loader = DataLoader(test, batch_size = batch_size, shuffle = False)
```

## CNN with 2 convolutional layer and 1 fully connected layer

```
In [13]: # Create CNN Model
class CNNModel(nn.Module):
    def __init__(self):
        super(CNNModel, self).__init__()

        self.cnn1 = nn.Conv2d(in_channels=3,out_channels=16,kernel_size=5,stroke=1,padding=0)
        self.relu1 = nn.ReLU()

        self.maxpool1 = nn.MaxPool2d(kernel_size=2)

        self.cnn2 = nn.Conv2d(in_channels=16,out_channels=32,kernel_size=5,stroke=1,padding=0)
        self.relu2 = nn.ReLU()

        self.maxpool2 = nn.MaxPool2d(kernel_size=2)

        self.fc1 = nn.Linear(15488,len(target_labels));

    def forward(self,x):

        out=self.cnn1(x)
        out=self.relu1(out)
        out=self.maxpool1(out)

        out=self.cnn2(out)
        out=self.relu2(out)
        out=self.maxpool2(out)

        out = out.view(out.size(0), -1)
        out=self.fc1(out)

        return out
```

```
In [14]: # Initialize Parameters and fit the model
model = CNNModel()
error = nn.CrossEntropyLoss()
learning_rate = 0.1
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

In [15]: *# CNN model training*

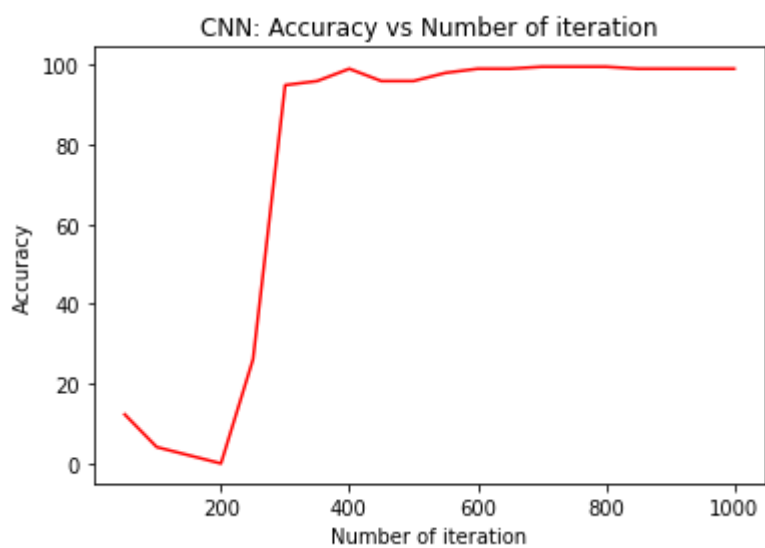

```
count = 0
loss_list = []
iteration_list = []
accuracy_list = []
for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        train = Variable(images.view(-1,3,100,100))
        #print(train.shape)
        labels = Variable(labels)
        optimizer.zero_grad()
        outputs = model(train)
        loss = error(outputs, labels)
        loss.backward()
        optimizer.step()
        count += 1

    if count % 50 == 0:
        correct = 0
        total = 0
        for images, labels in test_loader:
            test = Variable(images.view(-1,3,100,100))
            outputs = model(test)
            predicted = torch.max(outputs.data, 1)[1]
            total += len(labels)
            correct += (predicted == labels).sum()
        accuracy = 100 * correct / float(total)
        loss_list.append(loss.data)
        iteration_list.append(count)
        accuracy_list.append(accuracy)
        if count % 5 == 0:
            print('Iteration: {} Loss: {} Accuracy: {} %'.format(count
, loss.data, accuracy))
```

```
Iteration: 50 Loss: 3.9584412574768066 Accuracy: 12.307692527770996 %
Iteration: 100 Loss: 4.116231918334961 Accuracy: 4.102564334869385 %
Iteration: 150 Loss: 4.11642599105835 Accuracy: 2.0512821674346924 %
Iteration: 200 Loss: 4.173018455505371 Accuracy: 0.0 %
Iteration: 250 Loss: 3.294015884399414 Accuracy: 26.153846740722656 %
Iteration: 300 Loss: 0.19131530821323395 Accuracy: 94.87179565429688 %
Iteration: 350 Loss: 0.0065406630747020245 Accuracy: 95.8974380493164 %
Iteration: 400 Loss: 0.005745238158851862 Accuracy: 98.97435760498047 %
Iteration: 450 Loss: 0.0023304293863475323 Accuracy: 95.8974380493164 %
Iteration: 500 Loss: 0.02435666136443615 Accuracy: 95.8974380493164 %
Iteration: 550 Loss: 0.0016752103110775352 Accuracy: 97.94871520996094 %
Iteration: 600 Loss: 0.01584899052977562 Accuracy: 98.97435760498047 %
Iteration: 650 Loss: 0.0010343559551984072 Accuracy: 98.97435760498047 %
Iteration: 700 Loss: 0.013168991543352604 Accuracy: 99.4871826171875 %
Iteration: 750 Loss: 0.001453787088394165 Accuracy: 99.4871826171875 %
Iteration: 800 Loss: 0.0012704887194558978 Accuracy: 99.4871826171875 %
Iteration: 850 Loss: 0.00036563488538376987 Accuracy: 98.97435760498047 %
Iteration: 900 Loss: 0.005948456469923258 Accuracy: 98.97435760498047 %
Iteration: 950 Loss: 0.0012212992878630757 Accuracy: 98.97435760498047 %
Iteration: 1000 Loss: 0.0007654182845726609 Accuracy: 98.97435760498047 %
```

```
In [16]: # visualization loss
plt.plot(iteration_list,loss_list)
plt.xlabel("Number of iteration")
plt.ylabel("Loss")
plt.title("CNN: Loss vs Number of iteration")
plt.show()

# visualization accuracy
plt.plot(iteration_list,accuracy_list,color = "red")
plt.xlabel("Number of iteration")
plt.ylabel("Accuracy")
plt.title("CNN: Accuracy vs Number of iteration")
plt.show()
```



**CNN with 3 convolutional layer and 1 fully connected**

```

In [17]: # Create CNN Model
class CNNModel(nn.Module):
    def __init__(self):
        super(CNNModel, self).__init__()

        self.cnn1 = nn.Conv2d(in_channels=3,out_channels=16,kernel_size=5,stroke=1,padding=0)
        self.relu1 = nn.ReLU()

        self.maxpool1 = nn.MaxPool2d(kernel_size=2)

        self.cnn2 = nn.Conv2d(in_channels=16,out_channels=32,kernel_size=5,stroke=1,padding=0)
        self.relu2 = nn.ReLU()

        self.maxpool2 = nn.MaxPool2d(kernel_size=2)

        self.cnn3 = nn.Conv2d(in_channels=32,out_channels=64,kernel_size=5,stroke=1,padding=0)
        self.relu3 = nn.ReLU()

        self.maxpool3 = nn.MaxPool2d(kernel_size=2)

        self.fc1 = nn.Linear(5184,len(target_labels));

    def forward(self,x):

        out=self.cnn1(x)
        out=self.relu1(out)
        out=self.maxpool1(out)

        out=self.cnn2(out)
        out=self.relu2(out)
        out=self.maxpool2(out)

        out=self.cnn3(out)
        out=self.relu3(out)
        out=self.maxpool3(out)

        out = out.view(out.size(0), -1)
        out = self.fc1(out)

        return out

```

```

In [18]: # Initialize Parameters and fit the model
model = CNNModel()
error = nn.CrossEntropyLoss()
learning_rate = 0.02
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

```

In [19]: *# CNN model training*

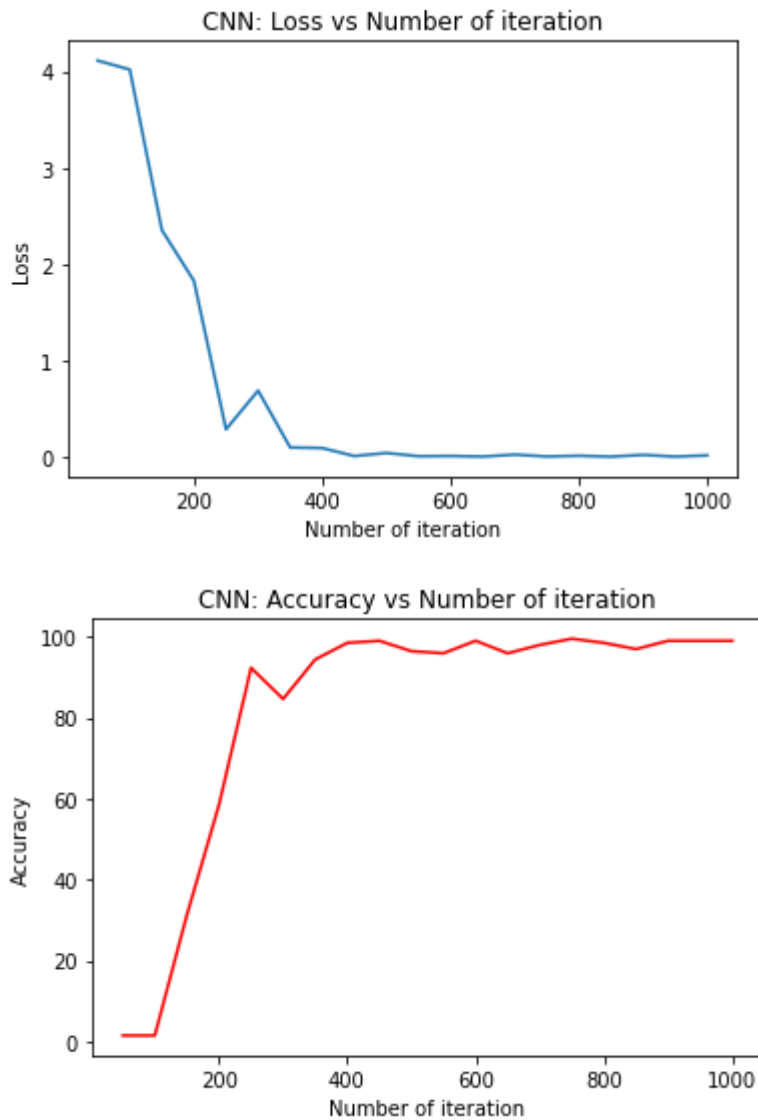
```
count = 0
loss_list = []
iteration_list = []
accuracy_list = []
for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        train = Variable(images.view(-1,3,100,100))
        labels = Variable(labels)
        optimizer.zero_grad()
        outputs = model(train)
        loss = error(outputs, labels)
        loss.backward()
        optimizer.step()
        count += 1

    if count % 50 == 0:
        correct = 0
        total = 0
        for images, labels in test_loader:
            test = Variable(images.view(-1,3,100,100))
            outputs = model(test)
            predicted = torch.max(outputs.data, 1)[1]
            total += len(labels)
            correct += (predicted == labels).sum()
        accuracy = 100 * correct / float(total)
        loss_list.append(loss.data)
        iteration_list.append(count)
        accuracy_list.append(accuracy)
        if count % 5 == 0:
            print('Iteration: {} Loss: {} Accuracy: {} %'.format(count
, loss.data, accuracy))
```

```
Iteration: 50 Loss: 4.111763000488281 Accuracy: 1.5384615659713745 %
Iteration: 100 Loss: 4.019026279449463 Accuracy: 1.5384615659713745 %
Iteration: 150 Loss: 2.353355646133423 Accuracy: 31.28205108642578 %
Iteration: 200 Loss: 1.8257973194122314 Accuracy: 58.46154022216797 %
Iteration: 250 Loss: 0.28569337725639343 Accuracy: 92.30769348144531 %
Iteration: 300 Loss: 0.6871101260185242 Accuracy: 84.61538696289062 %
Iteration: 350 Loss: 0.09699603170156479 Accuracy: 94.35897064208984 %
Iteration: 400 Loss: 0.09078714996576309 Accuracy: 98.46154022216797 %
Iteration: 450 Loss: 0.007882218807935715 Accuracy: 98.97435760498047 %
Iteration: 500 Loss: 0.03974886238574982 Accuracy: 96.4102554321289 %
Iteration: 550 Loss: 0.006276494357734919 Accuracy: 95.8974380493164 %
Iteration: 600 Loss: 0.008571005426347256 Accuracy: 98.97435760498047 %
Iteration: 650 Loss: 0.002602524124085903 Accuracy: 95.8974380493164 %
Iteration: 700 Loss: 0.022816266864538193 Accuracy: 97.94871520996094 %
Iteration: 750 Loss: 0.0035299521405249834 Accuracy: 99.4871826171875 %
Iteration: 800 Loss: 0.010895995423197746 Accuracy: 98.46154022216797 %
Iteration: 850 Loss: 0.001701981294900179 Accuracy: 96.92308044433594 %
Iteration: 900 Loss: 0.019954847171902657 Accuracy: 98.97435760498047 %
Iteration: 950 Loss: 0.0021584550850093365 Accuracy: 98.97435760498047 %
Iteration: 1000 Loss: 0.014821499586105347 Accuracy: 98.97435760498047 %
```

```
In [20]: # visualization loss
plt.plot(iteration_list,loss_list)
plt.xlabel("Number of iteration")
plt.ylabel("Loss")
plt.title("CNN: Loss vs Number of iteration")
plt.show()

# visualization accuracy
plt.plot(iteration_list,accuracy_list,color = "red")
plt.xlabel("Number of iteration")
plt.ylabel("Accuracy")
plt.title("CNN: Accuracy vs Number of iteration")
plt.show()
```



As we can see from above two model , if first I take 2 convolutional layer and 1 full connected layer with learning rate 0.1, next I take 3 convolutional layer and 1 fully connected layer with learning rate 0.02, keeping all parameters same I get good performance but If I take other learning rate with different numbers of convolutinal layer and fully connected layer then I get bad performance or underfitting model.