

# Weighted KNN

## Import dataset

```
In [18]: import sklearn
from sklearn.datasets import fetch_california_housing
# as_frame=True loads the data in a dataframe format, with other metadata besides i
california_housing = fetch_california_housing(as_frame=True)
# Select only the dataframe part and assign it to the df variable
df = california_housing.frame
```

```
In [19]: import pandas as pd
df.head()
```

```
Out[19]:
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	Med
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	

## Preprocessing Data for KNN Regression

```
In [20]: y = df['MedHouseVal']
X = df.drop(['MedHouseVal'], axis = 1)
```

```
In [21]: # .T transposes the results, transforming rows into columns
X.describe().T
```

```
Out[21]:
```

	count	mean	std	min	25%	50%	75%
MedInc	20640.0	3.870671	1.899822	0.499900	2.563400	3.534800	4.743250
HouseAge	20640.0	28.639486	12.585558	1.000000	18.000000	29.000000	37.000000
AveRooms	20640.0	5.429000	2.474173	0.846154	4.440716	5.229129	6.052380
AveBedrms	20640.0	1.096675	0.473911	0.333333	1.006079	1.048780	1.099520
Population	20640.0	1425.476744	1132.462122	3.000000	787.000000	1166.000000	1725.000000
AveOccup	20640.0	3.070655	10.386050	0.692308	2.429741	2.818116	3.282260
Latitude	20640.0	35.631861	2.135952	32.540000	33.930000	34.260000	37.710000
Longitude	20640.0	-119.569704	2.003532	-124.350000	-121.800000	-118.490000	-118.010000

## Splitting Data into Train and Test Sets

```
In [22]: from sklearn.model_selection import train_test_split

SEED = 42
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st
```

```
In [23]: print(len(X))      # 20640
print(len(X_train)) # 15480
print(len(X_test))  # 5160
```

```
20640
15480
5160
```

## Feature Scaling for KNN Regression

```
In [24]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
# Fit only on X_train
scaler.fit(X_train)

# Scale both X_train and X_test
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [25]: col_names=['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup',
scaled_df = pd.DataFrame(X_train, columns=col_names)
scaled_df.describe().T
```

```
Out[25]:
```

	count	mean	std	min	25%	50%	75%	max
<b>MedInc</b>	15480.0	2.172968e-16	1.000032	-1.774632	-0.688854	-0.175663	0.464450	5.842113
<b>HouseAge</b>	15480.0	-1.254954e-16	1.000032	-2.188261	-0.840224	0.032036	0.666407	1.855852
<b>AveRooms</b>	15480.0	-1.148163e-16	1.000032	-1.877586	-0.407008	-0.083940	0.257082	56.357392
<b>AveBedrms</b>	15480.0	1.239408e-16	1.000032	-1.740123	-0.205765	-0.108332	0.007435	55.925392
<b>Population</b>	15480.0	-7.874838e-17	1.000032	-1.246395	-0.558886	-0.227928	0.262056	29.971725
<b>AveOccup</b>	15480.0	2.672550e-17	1.000032	-0.201946	-0.056581	-0.024172	0.014501	103.737365
<b>Latitude</b>	15480.0	8.022581e-16	1.000032	-1.451215	-0.799820	-0.645172	0.971601	2.953905
<b>Longitude</b>	15480.0	2.169625e-15	1.000032	-2.380303	-1.106817	0.536231	0.785934	2.633738

## Training and Predicting KNN Regression

```
In [26]: from sklearn.neighbors import KNeighborsRegressor
regressor = KNeighborsRegressor(n_neighbors=5, weights="distance")
regressor.fit(X_train, y_train)
```

Out[26]: KNeighborsRegressor(weights='distance')

In [27]: `y_pred = regressor.predict(X_test)`

## Evaluating the Algorithm for KNN Regression

In [28]: `from sklearn.metrics import mean_absolute_error, mean_squared_error`

```
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False)

print(f'mae: {mae}')
print(f'mse: {mse}')
print(f'rmse: {rmse}')
```

```
mae: 0.44330658993325084
mse: 0.4284245302766481
rmse: 0.6545414656663457
```

In [29]: `regressor.score(X_test, y_test)`

Out[29]: 0.6762253110912666

In [30]: `y.describe()`

```
Out[30]: count    20640.000000
         mean       2.068558
         std       1.153956
         min       0.149990
         25%       1.196000
         50%       1.797000
         75%       2.647250
         max       5.000010
         Name: MedHouseVal, dtype: float64
```

## Finding the Best K for KNN Regression

In [31]: `error = []`


```
# Calculating MAE error for K values between 1 and 39
for i in range(1, 40):
    knn = KNeighborsRegressor(n_neighbors=i, weights="distance")
    knn.fit(X_train, y_train)
    pred_i = knn.predict(X_test)
    mae = mean_absolute_error(y_test, pred_i)
    error.append(mae)
```

In [32]: `import matplotlib.pyplot as plt`

```
plt.figure(figsize=(12, 6))
plt.plot(range(1, 40), error, color='red',
         linestyle='dashed', marker='o',
         markerfacecolor='blue', markersize=10)

plt.title('K Value MAE')
plt.xlabel('K Value')
plt.ylabel('Mean Absolute Error')
```

Out[32]: Text(0, 0.5, 'Mean Absolute Error')



```
In [33]: import numpy as np

print(min(error))
print(np.array(error).argmin())

0.43265872078512396
11
```

## KNN with 12 neighbours

```
In [34]: knn_reg12 = KNeighborsRegressor(n_neighbors=12, weights="distance")
knn_reg12.fit(X_train, y_train)
y_pred12 = knn_reg12.predict(X_test)
r2 = knn_reg12.score(X_test, y_test)

mae12 = mean_absolute_error(y_test, y_pred12)
mse12 = mean_squared_error(y_test, y_pred12)
rmse12 = mean_squared_error(y_test, y_pred12, squared=False)
print(f'r2: {r2}, \nmae: {mae12} \nmse: {mse12} \nrmse: {rmse12}')
```

r2: 0.6925746041555878,  
mae: 0.43265872078512396  
mse: 0.40679084969140783  
rmse: 0.6378015754852036

## Classification using K-Nearest Neighbors with Scikit-Learn

### Preprocessing Data for Classification

```
In [35]: # Creating 4 categories and assigning them to a MedHouseValCat column
df["MedHouseValCat"] = pd.qcut(df["MedHouseVal"], 4, retbins=False, labels=[1, 2, 3, 4])

In [36]: y = df['MedHouseValCat']
X = df.drop(['MedHouseVal', 'MedHouseValCat'], axis = 1)
```

## Splitting Data into Train and Test Sets

```
In [37]: from sklearn.model_selection import train_test_split

SEED = 42
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st
```

## Feature Scaling for Classification

```
In [38]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
scaler.fit(X_train)

X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

## Training and Predicting for Classification

```
In [39]: from sklearn.neighbors import KNeighborsClassifier

classifier = KNeighborsClassifier(weights="distance")
classifier.fit(X_train, y_train)
```

```
Out[39]: KNeighborsClassifier(weights='distance')
```

```
In [40]: y_pred = classifier.predict(X_test)
```

## Evaluating KNN for Classification

```
In [41]: acc = classifier.score(X_test, y_test)
print(acc)
```

```
0.6222868217054264
```


```
In [42]: from sklearn.metrics import classification_report, confusion_matrix
#importing Seaborn's to use the heatmap
import seaborn as sns

# Adding classes names for better interpretation
classes_names = ['class 1', 'class 2', 'class 3', 'class 4']
cm = pd.DataFrame(confusion_matrix(y_test, y_pred),
                  columns=classes_names, index = classes_names)

# Seaborn's heatmap to better visualize the confusion matrix
sns.heatmap(cm, annot=True, fmt='d');

print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
1	0.78	0.76	0.77	1292
2	0.51	0.54	0.52	1283
3	0.50	0.52	0.51	1292
4	0.72	0.66	0.69	1293
accuracy			0.62	5160
macro avg	0.63	0.62	0.62	5160
weighted avg	0.63	0.62	0.62	5160



## Finding the Best K for KNN Classification


```
In [43]: from sklearn.metrics import f1_score

f1s = []

# Calculating f1 score for K values between 1 and 40
for i in range(1, 40):
    knn = KNeighborsClassifier(n_neighbors=i, weights="distance")
    knn.fit(X_train, y_train)
    pred_i = knn.predict(X_test)
    # using average='weighted' to calculate a weighted average for the 4 classes
    f1s.append(f1_score(y_test, pred_i, average='weighted'))
```

```
In [44]: plt.figure(figsize=(12, 6))
plt.plot(range(1, 40), f1s, color='red', linestyle='dashed', marker='o',
         markerfacecolor='blue', markersize=10)
plt.title('F1 Score K Value')
plt.xlabel('K Value')
plt.ylabel('F1 Score')
```

```
Out[44]: Text(0, 0.5, 'F1 Score')
```



From the output, we can see that the f1-score is the highest when the value of the K is 10.

```
In [48]: classifier15 = KNeighborsClassifier(n_neighbors=10, weights="distance")
classifier15.fit(X_train, y_train)
y_pred15 = classifier15.predict(X_test)
print(classification_report(y_test, y_pred15))
```

	precision	recall	f1-score	support
1	0.79	0.77	0.78	1292
2	0.53	0.56	0.54	1283
3	0.51	0.55	0.53	1292
4	0.74	0.67	0.70	1293
accuracy			0.64	5160
macro avg	0.64	0.64	0.64	5160
weighted avg	0.64	0.64	0.64	5160

```
In [49]: acc = classifier.score(X_test, y_pred15)
print(acc)
```

```
0.8560077519379845
```

```
In [47]: cm = pd.DataFrame(confusion_matrix(y_test, y_pred15),
                           columns=classes_names, index = classes_names)

sns.heatmap(cm, annot=True, fmt='d');
```

